

CPSC 457

Assignment 1

Q1

a -

```
time python3 palindrome.py < t4.txt
```

Longest palindrome: redder

real 0m0.361s

user 0m0.317s

sys 0m0.006s

```
time ./slow-pali < t4.txt
```

Longest palindrome: redder

real 0m2.873s

user 0m1.402s

sys 0m1.457s

```
time python3 palindrome.py < t3.txt
```

Longest palindrome: __o.O.o__

real 0m0.028s

user 0m0.020s

sys 0m0.003s

```
time ./slow-pali < t3.txt
```

Longest palindrome: __o.O.o__

real 0m0.011s

user 0m0.005s

sys 0m0.000s

b-

The C++ code spent 0.000s running t3.txt and 1.402s running t4.txt in kernel mode compare to the python code spending 0.003s running t3.txt and 0.006s running t4.txt in kernel mode

c-

Usually C++ code is faster than python code because C++ is a compiled language while python is interpreted language which means there is a lot of overhead is needed to run the most basic python code. In this case, python is faster on longer files because the C++ code invokes a system call for every byte read which is very inefficient and causes a lot of time to be spent inside kernel mode in addition to the time spent context switching between user and kernel modes.

Q3- Fast-pali performance

```
time ./fast-pali < t3.txt
```

Longest palindrome: __o.O.o__

real 0m0.011s

user 0m0.003s

sys 0m0.003s

strace for fast-pali

```
time ./fast-pali < t4.txt
```

Longest palindrome: redder

real 0m0.109s

user 0m0.093s

sys 0m0.004s

```
strace -c time ./fast-pali < t4.txt
```

Longest palindrome: redder

```
0.08user 0.00system 0:00.08elapsed 97%CPU (0avgtext+0avgdata 4020maxresident)k
```

```
0inputs+0outputs (0major+382minor)pagefaults 0swaps
```

% time	seconds	usecs/call	calls	errors	syscall
--------	---------	------------	-------	--------	---------

76.95	0.003032	3032	1		wait4
9.37	0.000369	3	112		write
4.39	0.000173	173	1		execve
2.34	0.000092	5	18	16	openat
1.52	0.000060	8	7		mmap
1.19	0.000047	47	1		clone
0.91	0.000036	4	8	7	stat
0.69	0.000027	6	4		mprotect
0.63	0.000025	6	4		rt_sigaction
0.53	0.000021	21	1		munmap
0.38	0.000015	3	4		read
0.25	0.000010	3	3		Iseek
0.20	0.000008	4	2	1	arch_prctl
0.18	0.000007	3	2		close
0.18	0.000007	3	2		fstat
0.18	0.000007	7	1	1	access
0.10	0.000004	4	1		brk

100.00	0.003940	22	172	25	total
--------	----------	----	-----	----	-------

Strace for slow-pali

```

strace -c time ./slow-pali < t4.txt
Longest palindrome: redder
1.43user 1.45system 0:02.88elapsed 99%CPU (0avgtext+0avgdata 3048maxresident)k
48inputs+0outputs (0major+125minor)pagefaults 0swaps

% time   seconds usecs/call  calls  errors syscall
-----
99.96  1.451900  1451900    1      wait4
0.03  0.000445      3    112      write
0.00  0.000066     16      4      mprotect
0.00  0.000034     34      1      clone
0.00  0.000026      6      4      rt_sigaction
0.00  0.000024      3      7      mmap
0.00  0.000010      2      4      read
0.00  0.000008      8      1      munmap
0.00  0.000007      2      3      lseek
0.00  0.000004      0     18     16 openat
0.00  0.000003      1      2      close
0.00  0.000003      1      2      arch_prctl
0.00  0.000002      1      2      fstat
0.00  0.000000      0      8      7 stat
0.00  0.000000      0      1      brk
0.00  0.000000      0      1      access
0.00  0.000000      0      1      execve
-----
100.00  1.452532   8444    172     25 total

```

a- By comparing the time it took fast-pali to process t3.txt and t4.txt, we can see that fast-pali is slightly faster than slow-pali on the small t3.txt file and much faster than slow-pali on the large t4.txt. Fast-pali is faster than slow-pali because fast pali preforms far less system calls compared to slow-pali. As we can by the strace output above fast-pali makes 22 system calls compared to the 8444 system calls done by slow-pali.

b- By comparing the time it took fast-pali to process t3.txt and t4.txt, we can see that fast-pali is slightly faster than palindrome.py on the small t3.txt file and much faster than it on the large t4.txt. fast-pali is simply faster due to interpreter overhead needed to run the python code.